

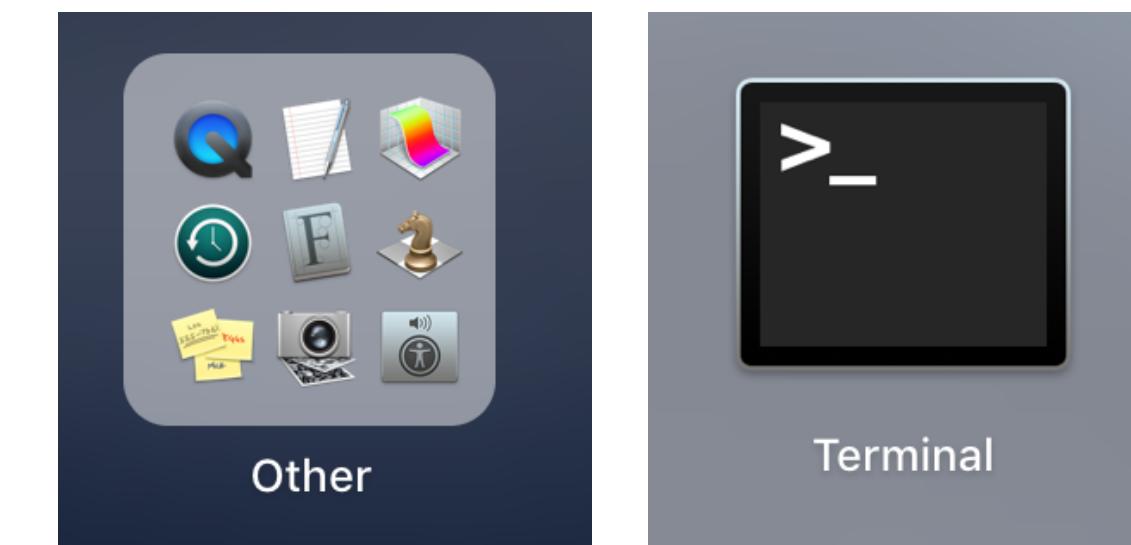
2/18/2020 Week 6 Module 1

Running molecular dynamics on a remote computer

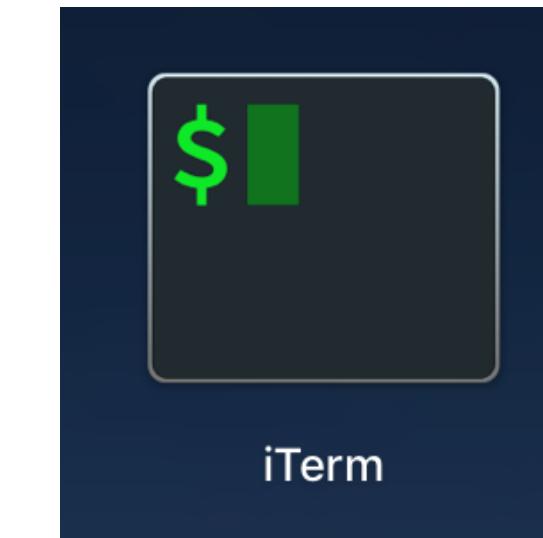
- Useful molecular dynamics simulations and other scientific calculations often require too much computer power to run on a local laptop or desktop
- In this module, you will learn how to
 - log into a remote computer
 - install OpenMM on the remote computer
 - transfer data to and from a remote computer
 - submit a calculation to a queue on a computing cluster
- We will use remote computers belonging to XSEDE, the “Extreme Science and Engineering Discovery Environment”, an organization that manages computing resources for open science throughout the United States

Review: The UNIX Terminal

- Predecessor and very similar to the terminal used in
 - LINUX - many scientific computing clusters
 - Mac OS X - this lab and many of your own laptops/computers
- Useful for
 - accessing programs that run with a command line interface (CLI)
 - automation



- To start the terminal on a Mac:

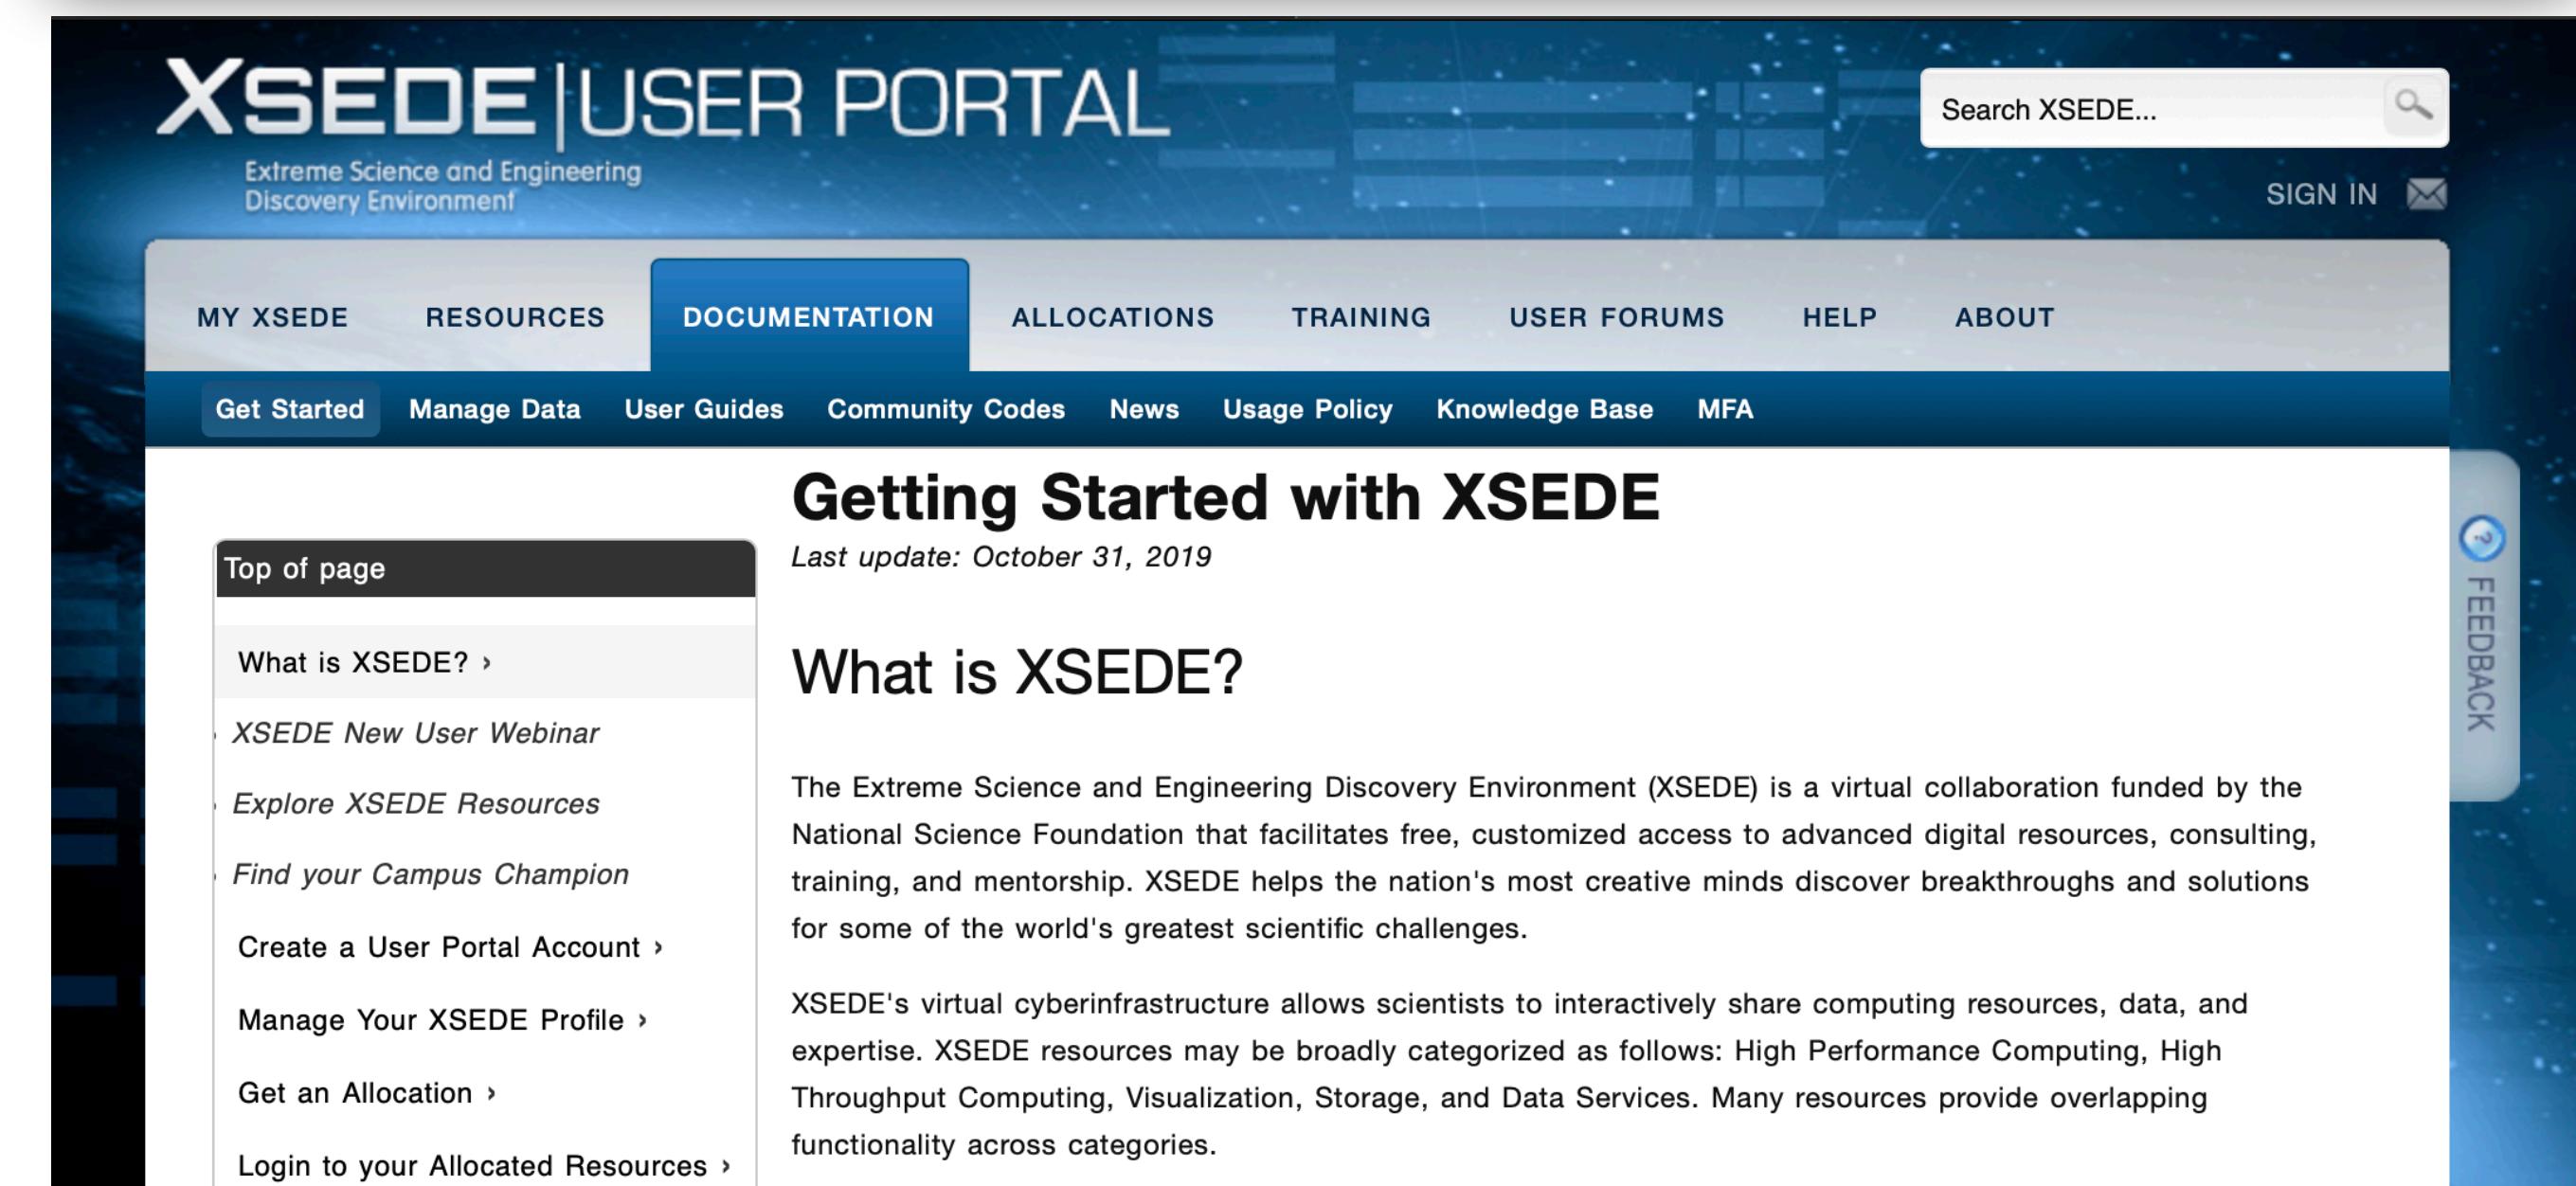
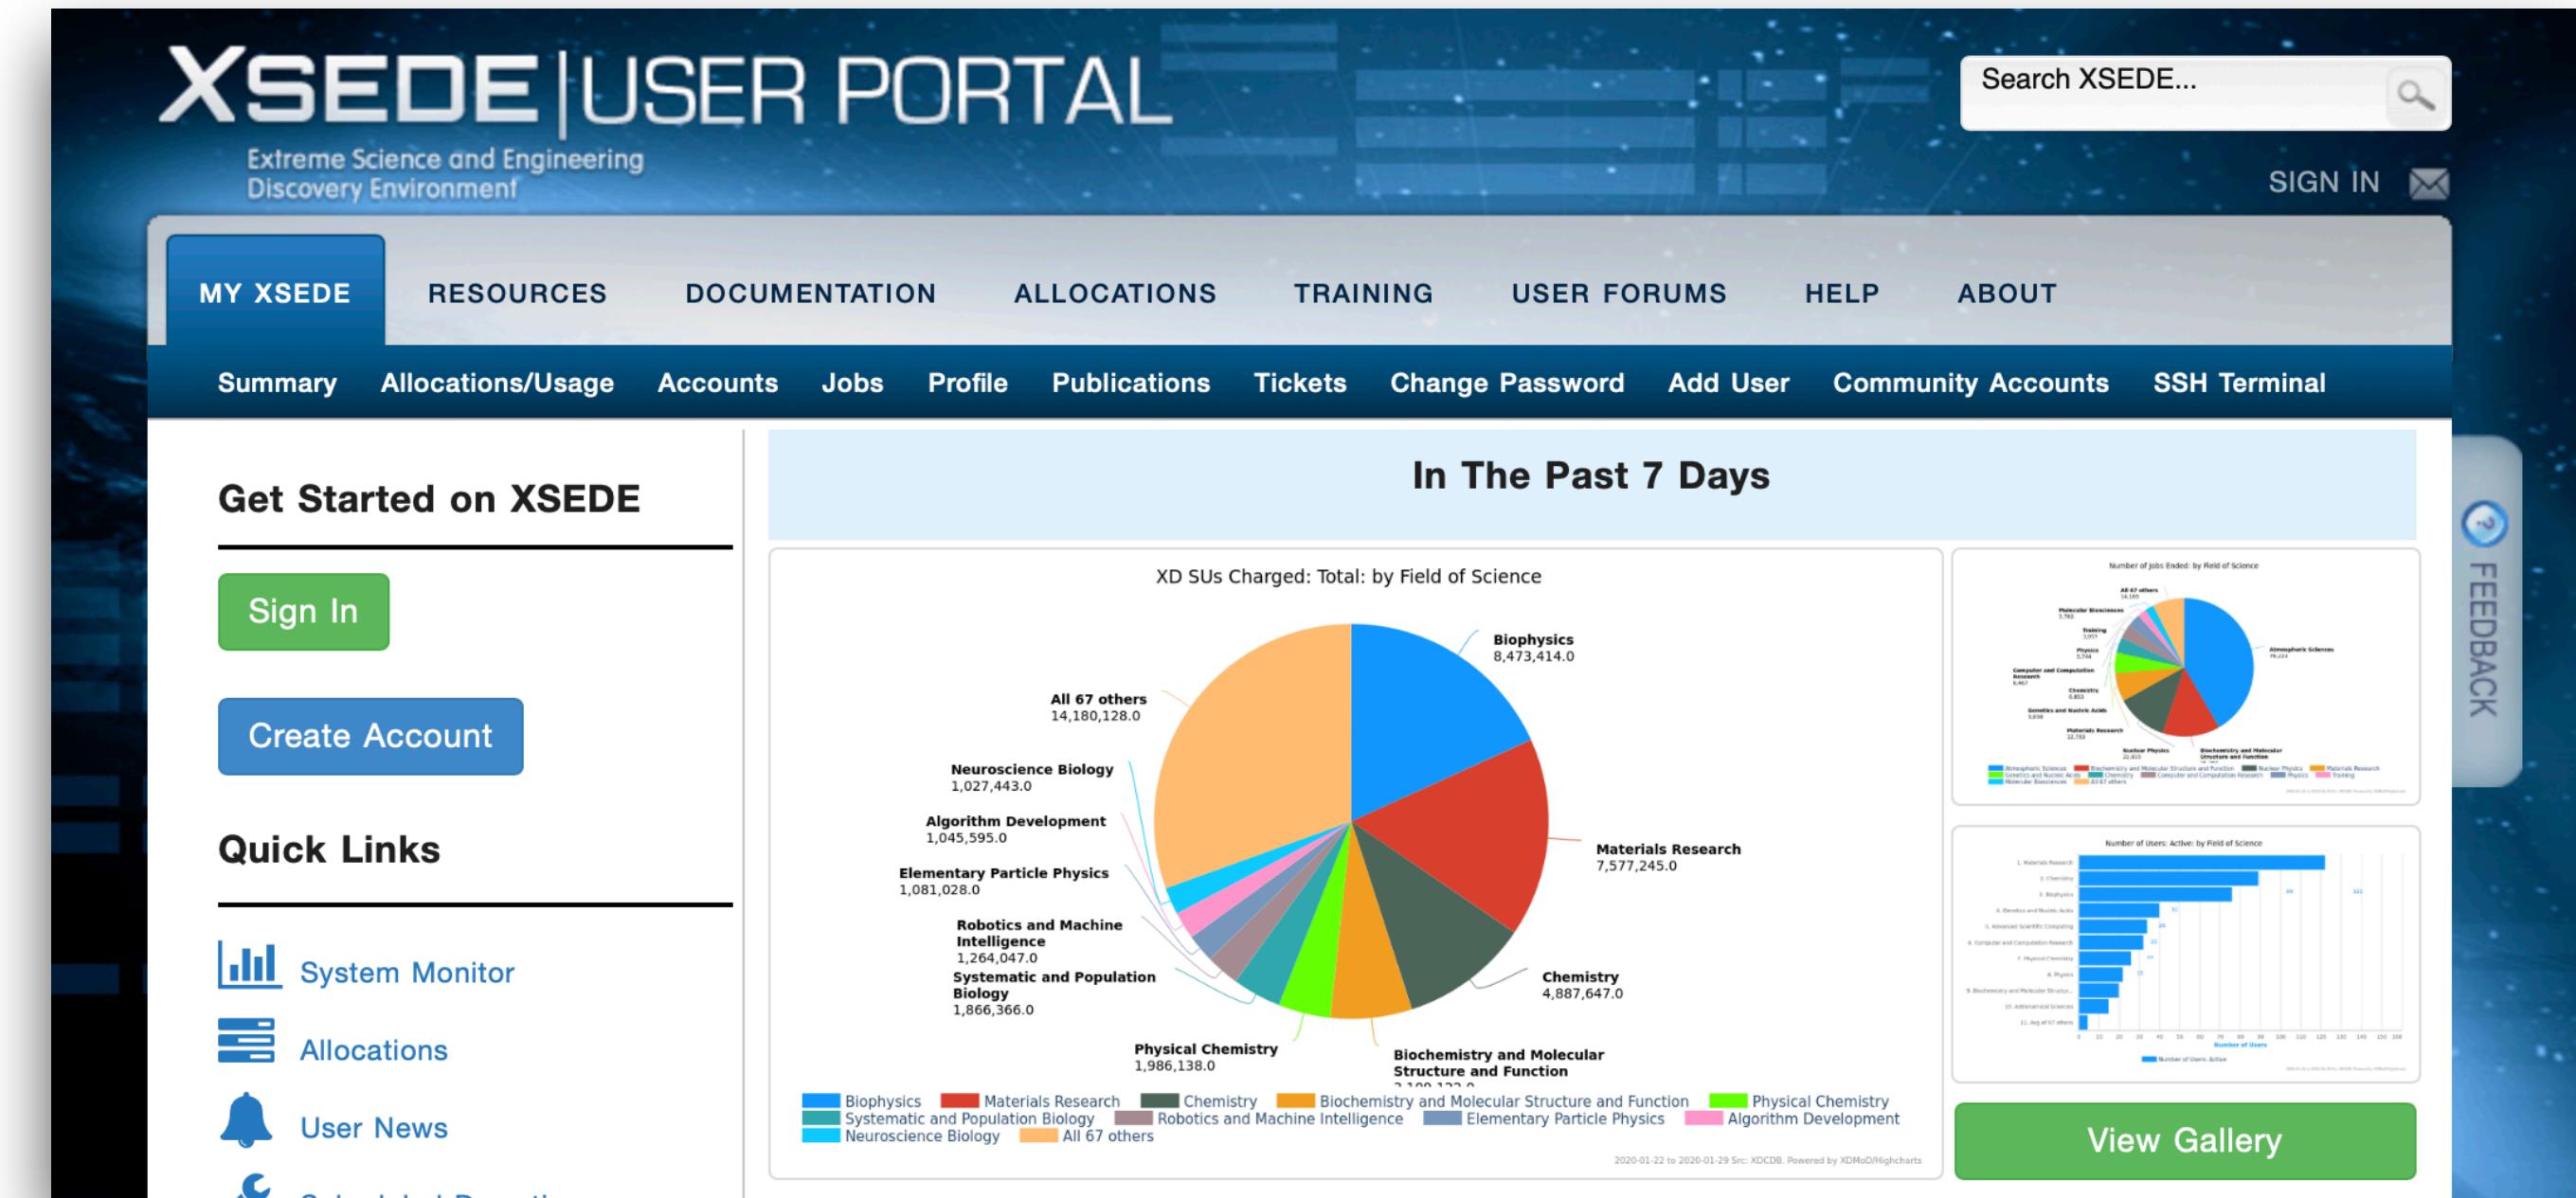


- I actually usually use another program:

Review: UNIX Terminal Essentials

- Try the following commands
 - echo \$SHELL - echo reports the value of a variable. \$SHELL is the terminal interface you are using, and will affect the details of all other commands. bash is a popular shell.
 - File operations
 - ls – list files and directories
 - cp – copy files
 - rm – remove files and directories
 - mv – rename or move files and directories to another location
 - Directory management
 - pwd – confirm current directory
 - cd – change directory
 - mkdir – make new directory
 - rmdir – remove directory
- Also see
 - <https://www.unixtutorial.org/basic-unix-commands>
 - An Introduction to Linux (<https://www.youtube.com/watch?v=IVquJh3DXUA>). Work with the terminal starts at 3:35.

- A lot of what we will work through today is covered in the “Getting started with XSEDE” tutorial.
- This can be found at <http://portal.xsede.org/>. Click on the tabs marked “Documentation” and “Get Started”.
- You should already have
 - an XSEDE account and access to the class allocation
 - signed up for Multi-Factor Authentication with Duo (<https://portal.xsede.org/mfa>)
- So you can start from “Login to your Allocated Resources” on the left pane.



Cluster computing



On 7/22/2019, your professor visited the Texas Advanced Computing Center (TACC), home of an XSEDE computing cluster, Stampede2.

- The XSEDE supercomputers are not especially fast computers, but are computing clusters, groups of computing nodes that work together
- The most common way for computational scientists to utilize these computers is to submit a batch job to the queue
 - a batch job contains
 - information about the job, e.g.
 - allocation
 - maximum duration
 - type and amount of resource, e.g. CPUs or GPUs
 - a series of commands, like if you typed them into the terminal
 - The queue
 - is usually managed by a master node
 - prioritizes jobs
 - distributes jobs to worker nodes

Logging In

- We will follow the “Single Sign On via SSH Login Hub”, so you can enter
 - ssh username@login.xsede.org
 - your XUP passcode
 - then ask for a push to your phone
 - more information is at <https://portal.xsede.org/web/xup/single-sign-on-hub>
- From the login portal, you can log into bridges/comet by entering
 - `gsissh bridges' or
 - `gsissh comet'
- To log out of a remote computer, press control-D

```
Minh-IIT-MBP2018:[~]: ssh daveminh@login.xsede.org
Please login to this system using your XSEDE username and password:
password:
Duo two-factor login for daveminh

Enter a passcode or select one of the following options:

 1. Duo Push to XXX-XXX-4139
 2. Phone call to XXX-XXX-4139

Passcode or option (1-2): 1
Success. Logging you in...
Last login: Sat Feb  1 13:41:10 2020 from 207.237.205.39

# Welcome to the XSEDE Single Sign-On (SSO) Hub!
#
# This system is for use by authorized users only, and is subject to the XSEDE
# Acceptable Use Policy, described at https://www.xsede.org/usage-policies.
# All activities on this system may be monitored and logged.
#
# Your storage on this system is limited to 100MB. Backup is not provided.
#
# From this system, you may login to other XSEDE system login hosts on which
# you currently have an active account. To see a list of your accounts, visit:
# https://portal.xsede.org/group/xup/accounts
#
# To login to an XSEDE system login host, enter: gsissh <login-host>
# where <login-host> is the hostname, alias or IP address of the login host.
# The following default gsissh host aliases have been defined:
#
#      bridges  comet  mason  osg  rmacc-summit  stampede
#                  stampede2  supermic  wrangler-iu  wrangler-tacc
#
# For example, to login to the Comet system at SDSC, enter: gsissh comet
#
# E-mail help@xsede.org if you require assistance in the use of this system.
```

```
[daveminh@ssohub ~]$ gsissh comet
Last login: Sat Feb  1 10:46:54 2020 from 149.165.168.51
Rocks 7.0 (Manzanita)
Profile built 13:03 03-Dec-2019
```

```
Kickstarted 14:18 03-Dec-2019
```

WELCOME TO



```
comet-ln3:[~]:
```

The Bridges File System

- As described in the [Bridges user guide](#) under “File Spaces”, Bridges allows you to store data in
 - your home directory (\$HOME) - shared across nodes, backed up, smaller quota (10 GB)
 - a “scratch” directory (\$SCRATCH) - shared across nodes, not backed up, faster than \$HOME, larger quota shared by class
 - the worker node for a job (\$LOCAL) - only accessible by local node, not backed up and immediately deleted after job, faster than shared disk space, many terabytes of available space
 - the local memory for a job (\$RAMDISK) - only accessible by local node, not backed up and immediately deleted after job, fastest storage
 - \$(VARIABLE_NAME) is a variable that contains the actual location. The variable is available while logged onto Bridges
- Try `cd \$SCRATCH`, ‘pwd’, `cd \$HOME`, and ‘pwd’

Installing OpenMM

- The easiest way to install OpenMM is through the conda package manager
- Conda is available on bridges, but needs to be loaded by entering `module load anaconda`
- Then you need to create an conda environment that you can write to by entering `conda create --name openmm`

```
br005:[~]: conda create --name openmm
Collecting package metadata (current_repodata.json): done
Solving environment: done

==> WARNING: A newer version of conda exists. <==
    current version: 4.7.12
    latest version: 4.8.1

Please update conda by running

$ conda update -n base -c defaults conda

## Package Plan ##

environment location: /home/dminh/.conda/envs/openmm

Proceed ([y]/n)? y

Preparing transaction: done
Verifying transaction: done
Executing transaction: done
#
# To activate this environment, use
#
#     $ conda activate openmm
#
# To deactivate an active environment, use
#
#     $ conda deactivate
```

- To allow you to activate the environment, enter `conda init bash` and then `bash`
- Then to activate the environment, enter `conda activate openmm`
- Finally, to install OpenMM, enter
 - `conda install -c omnia -c conda-forge openmm`
 - `conda init bash`
- OpenMM may be tested using `python -m simtk.testInstallation`

```
(openmm) br006:[~]: conda install -c omnia -c conda-forge openmm
Collecting package metadata (current_repotdata.json): done
Solving environment: done
```

```
==> WARNING: A newer version of conda exists. <==
current version: 4.7.12
latest version: 4.8.1
```

Please update conda by running

```
$ conda update -n base -c defaults conda
```

Package Plan

environment location: /home/dminh/.conda/envs/openmm

added / updated specs:

- openmm

The following packages will be downloaded:

package	build		
-----	-----		
_libgcc_mutex-0.1	conda_forge	3 KB	conda-forge
_openmp_mutex-4.5	0_gnu	435 KB	conda-forge
ca-certificates-2019.11.28	hecc5488_0	145 KB	conda-forge
certifi-2019.11.28	py37_0	148 KB	conda-forge
cython-0.29.14	py37he1b5a44_0	2.2 MB	conda-forge
ld_impl_linux-64-2.33.1	h53a641e_8	589 KB	conda-forge
libblas-3.8.0	14_openblas	10 KB	conda-forge

Transferring Data

- After logging into the XSEDE Single Sign-On Hub, you can transfer data
 - using various commands in the terminal, including rsync, scp, and sftp
 - using specialized programs, as specified in the user manuals, but this requires installing programs, which you might not want to do
- On the terminal, I recommend rsync because you don't need to copy (and wait for) the whole directory
- A good template is: rsync -Cuavz --port 2222 source_path destination_path
 - Cuavz are options, which can be seen by `man rsync`
 - port 2222 means that you go through the XSEDE SSO
 - the paths have the format user_name@directory/file_name
 - if user_name@ is omitted, it assumes the same user name as on the current system
 - if directory is omitted, the home directory is assumed
 - if file_name is omitted, the whole directory is assumed

- Try copying a local directory with the ubiquitin pdb file and python script for OpenMM

```
Minh-IIT-MBP2018:[~/Documents/GitHub/Chem456/static_files/tutorials]: rsync -Cua
vz --port 2222 ubq-md dminh@data.bridges.psc.edu:/home/dminh/
building file list ... done
ubq-md/
ubq-md/1ubq.pdb
ubq-md/MD_ubq.py

sent 14309 bytes received 70 bytes 9586.00 bytes/sec
total size is 52573 speedup is 3.66
Minh-IIT-MBP2018:[~/Documents/GitHub/Chem456/static_files/tutorials]:
```

If you don't have them on your local computer, you can download them from github:
https://github.com/daveminh/Chem456/tree/master/static_files/tutorials/ubq-md

- Once you have installed OpenMM and copied the data, you can simply run the python script on the login terminal
- However, this is not the way that you're supposed to do things
- If you try to run a big calculation on the login terminal the system administrators will get mad at you!

```
(openmm) br006:[~]: python -m simtk.testInstallation
OpenMM Version: 7.4.1
Git Revision: 068f120206160d5151c9af0baf810384bba8d052

There are 2 Platforms available:

1 Reference - Successfully computed forces
2 CPU - Successfully computed forces

Median difference in forces between platforms:

Reference vs. CPU: 6.30481e-06

All differences are within tolerance.
(openmm) br006:[~]: ls
scripts software ubq-md
(openmm) br006:[~]: cd ubq-md/
(openmm) br006:[~/ubq-md]: ls
1ubq.pdb MD_ubq.py
(openmm) br006:[~/ubq-md]: python MD_ubq.py
Minimizing...
Running Production...
##"Progress (%)" "Step" "Potential Energy (kJ/mole)" "Temperature (K)"
"Speed (ns/day)" "Time Remaining"
10.0% 100 -12937.54104464231 178.8691182900329 0 --
20.0% 200 -12786.424474924024 206.88709731573758 21.9 0:06
30.0% 300 -12425.520220420438 214.78270910359848 21.8 0:05
40.0% 400 -12333.981042607265 242.67256593500656 21.8 0:04
50.0% 500 -12170.023875048062 252.83736423813997 21.8 0:03
60.0% 600 -12043.26256936375 262.77890676191413 21.7 0:03
70.0% 700 -12022.84441006196 274.6817644781113 21.8 0:02
80.0% 800 -11904.676295196514 279.1694256963495 21.8 0:01
90.0% 900 -11873.317454425185 274.09491413581435 21.7 0:00
100.0% 1000 -11727.128648981274 275.21880048454915 21.7 0:00
Done!
```

Submitting jobs to Bridges

- As described in the [Bridges user guide](#) under “Running Jobs”, you can submit a job to the cluster using a batch script
- Bridges uses the SLURM (Simple Linux Utility for Resource Management) scheduler
- A batch job consists of
 - a line that describes the shell, e.g. `#!/bin/bash`
 - SLURM directives that describe job properties, starting with `#SBATCH`
 - many directives can be command-line arguments, but it is generally better to save job properties into the job script
 - other schedulers will have different directives
 - executable commands, details of which will depend on the shell
 - optionally, comments after `#`

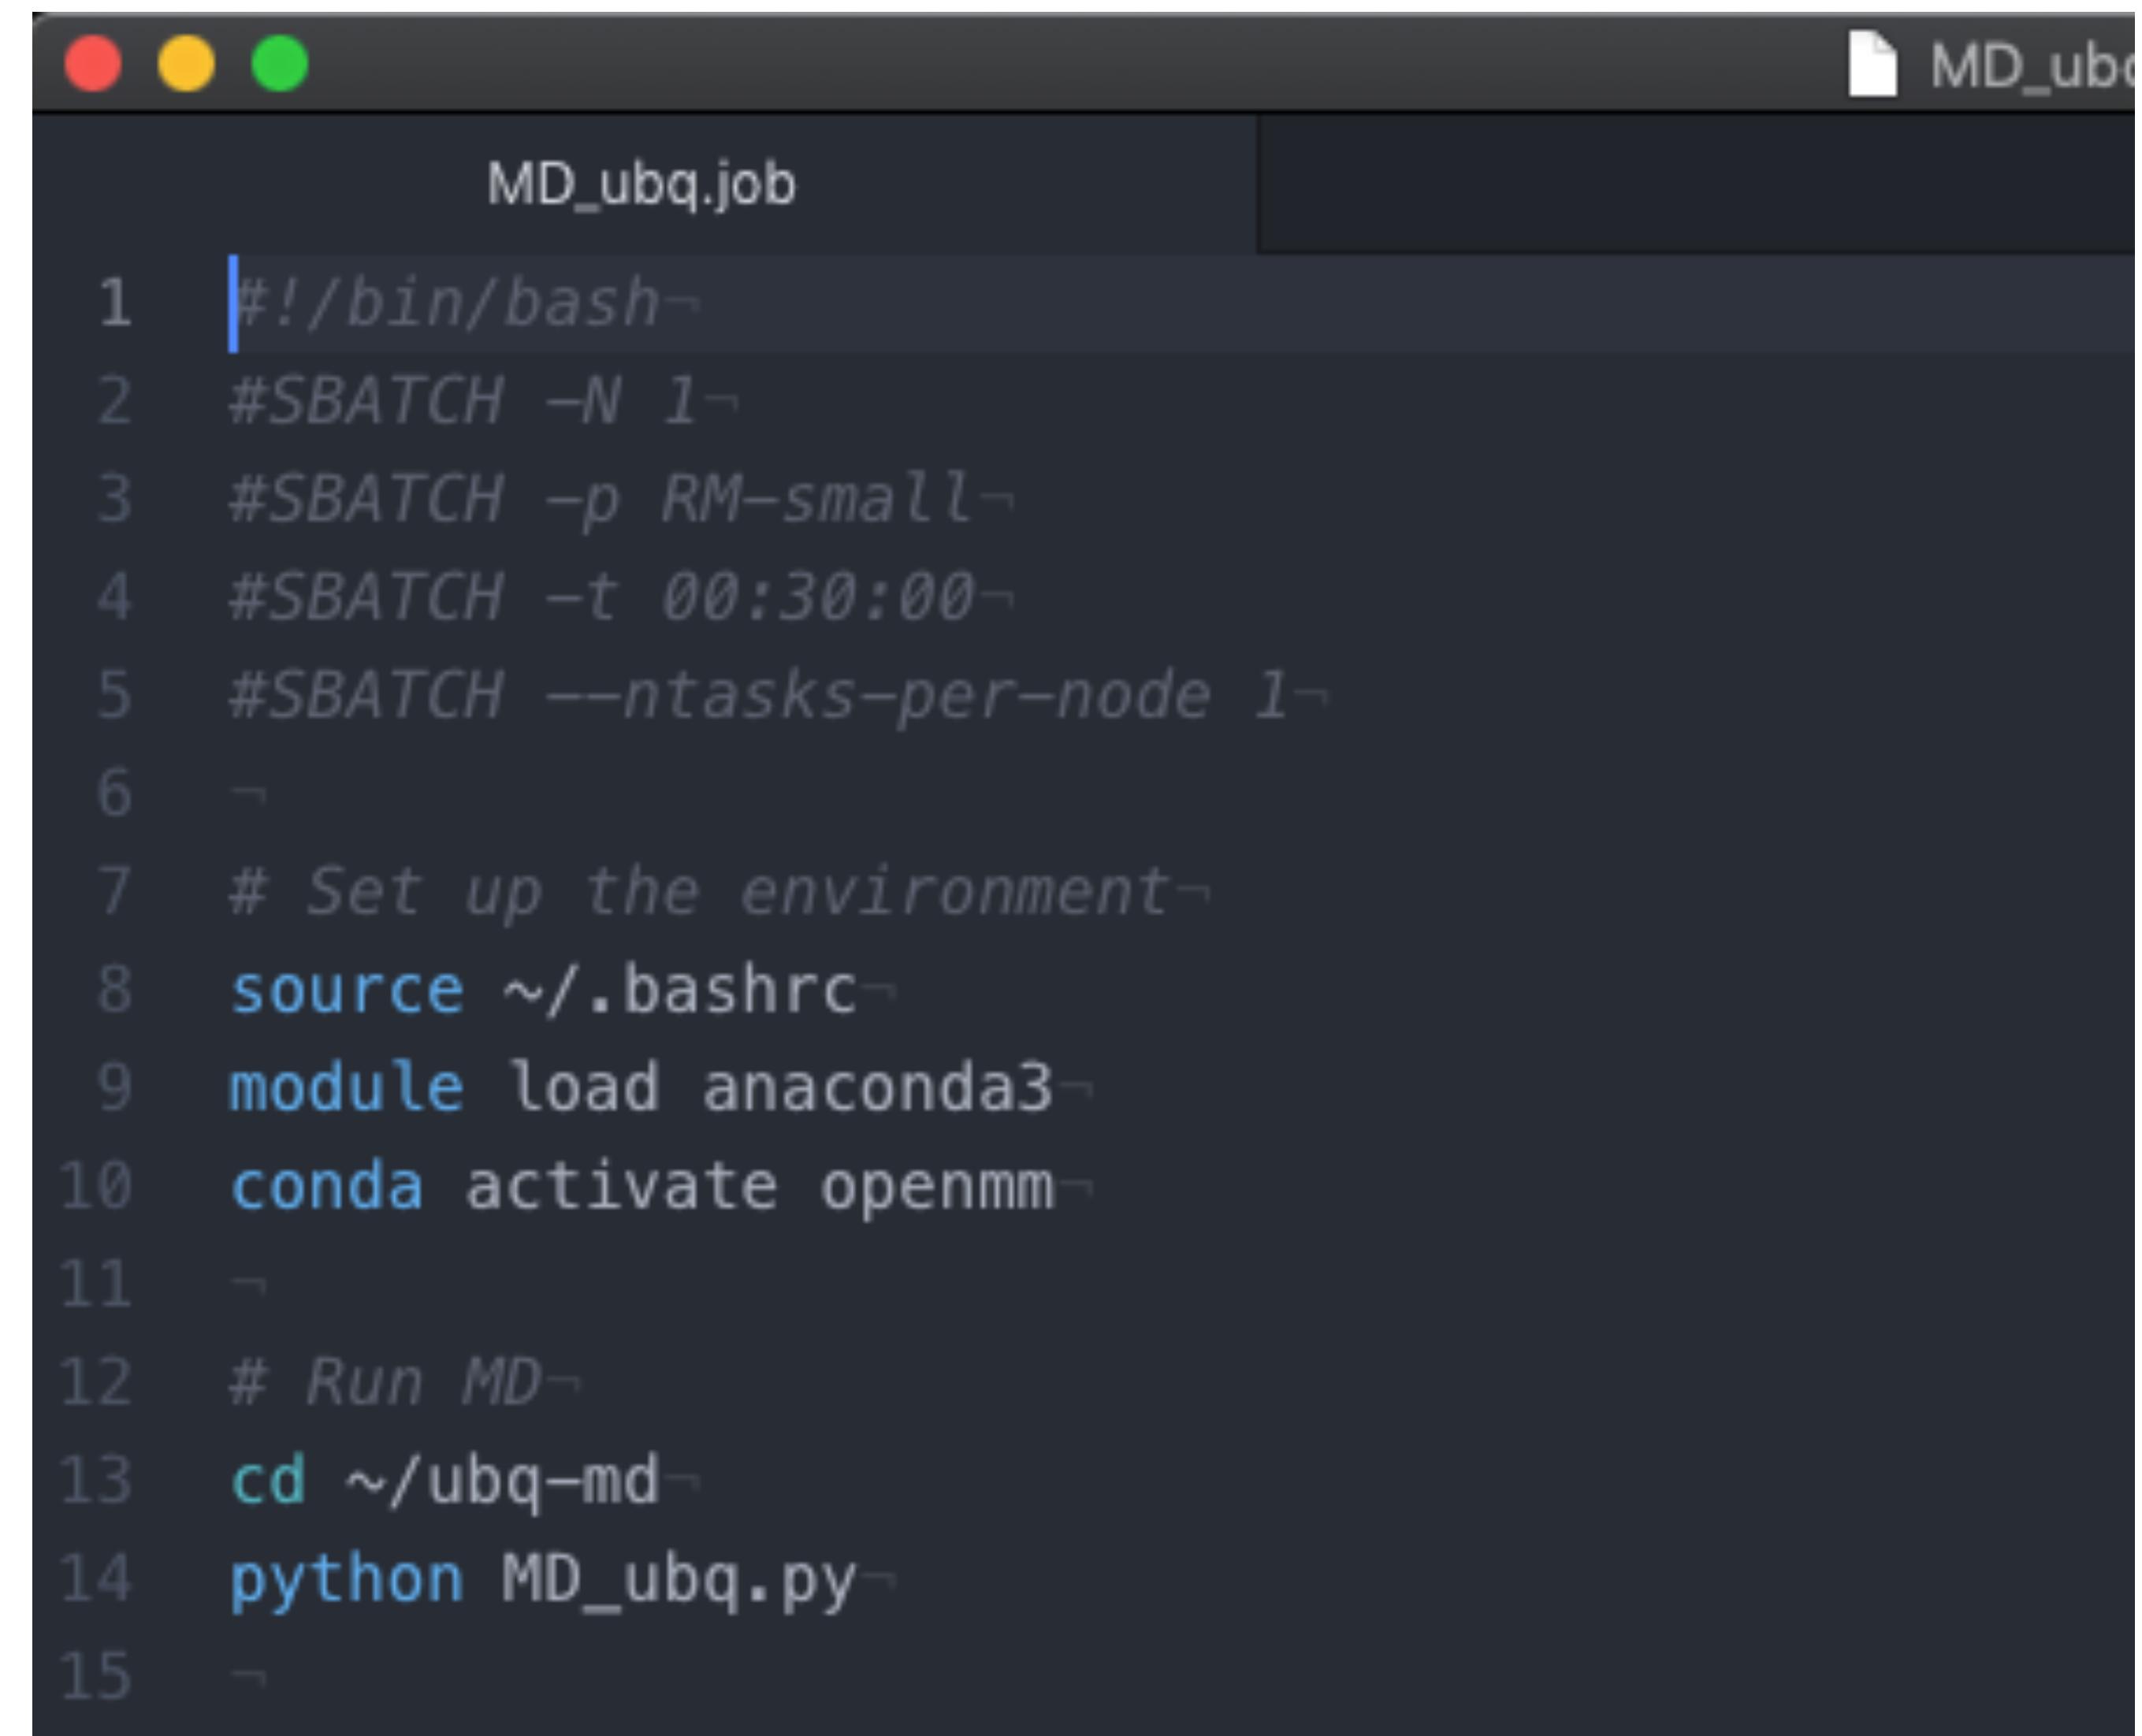
- SLURM has its own commands (<https://slurm.schedmd.com/quickstart.html>)
- Let's try two simple ones
 - sinfo - information about the Bridges partitions. We want “RM-small” for our first job.
Shorter jobs can be squeezed in between longer jobs.
 - squeue - lists jobs currently on the queue

```

● ● ● daveminh@ssohub:~ 7%1
br005:[~/ubq-md]: ls
1ubq.pdb MD_ubq.py trajectory.dcd ubq_mod.pdb
br005:[~/ubq-md]: sinfo
PARTITION AVAIL TIMELIMIT NODES STATE NODELIST
RM*       up 3-00:00:00   1 drain* r389
RM*       up 3-00:00:00   3 comp r[076,410,513]
RM*       up 3-00:00:00   5 resv r[212,383,385-386,388]
RM*       up 3-00:00:00  44 mix r[562,565,571-572,578-579,584-585,587,
590,596,599-600,603-604,613-614,618,620,622,624,626-628,636,641,656,662,674,687-
688,690-691,693,697,701-703,714-715,723-724,729,741]
RM*       up 3-00:00:00  678 alloc r[006-075,077-211,213-235,237-239,241-
277,279-309,311-382,384,387,390-409,411-437,439-490,492-512,514-561,563,566-570,
573-574,576-577,580-583,586,588-589,591-595,597-598,601-602,605-612,615-617,619,
621,623,625,629-635,637-640,642-655,657-661,663-673,675-686,689,692,694-696,698-
700,704-713,716-722,725-728,730-740,742-744]
RM*       up 3-00:00:00   8 idle r[236,240,278,310,438,491,564,575]
RM-shared  up 3-00:00:00  46 mix r[562,565,571-572,578-579,584-585,587,
590,596,599-600,603-604,613-614,618,620,622,624,626-628,636,641,656,662,674,687-
688,690-691,693,697,701-703,714-715,723-724,729,741,751-752]
RM-shared  up 3-00:00:00  152 alloc r[553-561,563,566-570,573-574,576-577,
580-583,586,588-589,591-595,597-598,601-602,605-612,615-617,619,621,623,625,629-
635,637-640,642-655,657-661,663-673,675-686,689,692,694-696,698-700,704-713,716-
722,725-728,730-740,742-750]
RM-shared  up 3-00:00:00   2 idle r[564,575]
RM-small   up 8:00:00    1 mix r005
RM-small   up 8:00:00    3 alloc r[001-003]
RM-small   up 8:00:00    1 idle r004
GPU        up 2-00:00:00  2 mix gpu[003,026]
GPU        up 2-00:00:00  25 alloc gpu[001-002,004-007,009,012,017-025,03
5-037,039-040,042-044]
GPU        up 2-00:00:00  17 idle gpu[008,010-011,013-016,027-034,038,04
1]
GPU-shared  up 2-00:00:00  2 mix gpu[003,026]
GPU-shared  up 2-00:00:00  25 alloc gpu[001-002,004-007,009,012,017-025,03
5-037,039-040,042-044]
GPU-shared  up 2-00:00:00  17 idle gpu[008,010-011,013-016,027-034,038,04
1]
GPU-small  up 8:00:00    2 mix gpu[045-046]
GPU-small  up 8:00:00    2 idle gpu[047-048]
GPU-AI     up 2-00:00:00  6 mix gpu[051-054,057-058]
GPU-AI     up 2-00:00:00  4 alloc gpu[049-050,055-056]
LM         up 14-00:00:0 38 mix l[001-003,006,008-012,014-020,022-028,
030-041],xl[002-004]
LM         up 14-00:00:0  8 alloc l[004-005,007,013,021,029,042],xl001
XLM        up 14-00:00:0  3 mix xl[002-004]
XLM        up 14-00:00:0  1 alloc xl001
DBMI       up 2-00:00:00  4 mix dr[001-004]
DBMI       up 2-00:00:00  4 idle dr[005-008]
DBMI-GPU   up 2-00:00:00  1 mix dgpu003
DBMI-GPU   up 2-00:00:00  2 alloc dgpu[002,004]
DBMI-GPU   up 2-00:00:00  1 idle dgpu001
br005:[~/ubq-md]: 

```

- Let's write a batch script called MD_ubq.job and put it in the ~/ubq-md/ directory. You can
 - use a terminal text editor like `vi'
 - create it on your local computer and transfer it to bridges



A screenshot of a terminal window titled "MD_ubq.job". The window shows a bash script with the following content:

```
1 #!/bin/bash-
2 #SBATCH -N 1-
3 #SBATCH -p RM-small-
4 #SBATCH -t 00:30:00-
5 #SBATCH --ntasks-per-node 1-
6 -
7 # Set up the environment-
8 source ~/.bashrc-
9 module load anaconda3-
10 conda activate openmm-
11 -
12 # Run MD-
13 cd ~/ubq-md-
14 python MD_ubq.py-
15 -
```

- Once you have created the job file, you can submit it to the queue using `sbatch MD_ubq.job'
- After submitting the job you can log out or be disconnected from the server and it will still run
- You can use `squeue -u *user_name*' to check the status of your job(s)
- The calculation will complete the same way as an interactive job
- The output that would be sent to the terminal in an interactive job is sent to a file

```
br006:[~/ubq-md]: sbatch MD_ubq.job
Submitted batch job 7589729
br006:[~/ubq-md]: squeue -u dminh
             JOBID PARTITION     NAME   USER ST      TIME  NODES NODELIST(REA
SON)
              7589729  RM-small  MD_ubq.j    dminh R       0:06      1 r005
```

```
br006:[~/ubq-md]: ls -ltr
total 312
-rw-r--r-- 1 dminh mc4s8bp 50895 Jan 31 18:44 1ubq.pdb
-rw-r--r-- 1 dminh mc4s8bp 1678 Jan 31 18:50 MD_ubq.py
-rw-r--r-- 1 dminh mc4s8bp 221 Feb  5 17:17 MD_ubq.job
-rw-r--r-- 1 dminh mc4s8bp 99863 Feb  5 17:23 ubq_mod.pdb
-rw-r--r-- 1 dminh mc4s8bp 148796 Feb  5 17:23 trajectory.dcd
-rw-r--r-- 1 dminh mc4s8bp 727 Feb  5 17:23 slurm-7589729.out
br006:[~/ubq-md]: more slurm-7589729.out
Minimizing...
Running Production...
##"Progress (%)" "Step"  "Potential Energy (kJ/mole)"    "Temperature (K)"
##"Speed (ns/day)"        "Time Remaining"
10.0%  100    -12891.608289052907  182.75168707244657  0      --
20.0%  200    -12679.091283728229  208.12074466763545  8.22   0:16
30.0%  300    -12544.950513910171  222.9911557564418   8.27   0:14
40.0%  400    -12417.294668972278  241.63949664412493  8.29   0:12
50.0%  500    -12182.148295443869  250.84879213741047  8.31   0:10
60.0%  600    -12086.427954357983  250.25410151752132  8.31   0:08
70.0%  700    -12201.427649847661  275.4539256800823   8.3    0:06
80.0%  800    -12071.165875336525  278.7590849918891   8.29   0:04
90.0%  900    -11779.891281989974  278.2551306095054   8.28   0:02
100.0% 1000   -11804.03142004821   289.69894096208446  8.27   0:00
Done!
```

Some Complications

- For your project, you will want different options in your scripts to
 - run MD on GPUs
 - run calculation for longer
 - not overwrite previous data
 - store data on \$LOCAL or \$SCRATCH